

## Business case Introduction

**Cars24** is one of the leading used-car buying and selling platform in India.

Current methodology to estimate the selling price has three major components

1. Past records of the "Selling Prices" of the same *Make* and *Model*
2. Rule based calculations for the depreciation of the car
3. Rule based calculation for condition of the vehicle

As the **Pricing Data Scientist working at Cars24**, you are given a task

GOAL: To automate the process of estimating price of the used car

### ▼ Exploratory Data Analysis

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
!gdown 18QuoyuqmNKmR9mba5he4ojPc4KzkIs3t
```

```
Downloading...
```

```
From: https://drive.google.com/uc?id=18QuoyuqmNKmR9mba5he4ojPc4KzkIs3t
```

```
To: /content/train-cars24-car-price.csv
```

```
100% 1.25M/1.25M [00:00<00:00, 119MB/s]
```

### ▼ Overview

```
df = pd.read_csv("train-cars24-car-price.csv")
df.head()
```

	full_name	selling_price	year	seller_type	km_driven	fuel_type	transmission_type	mileage	enq
0	Maruti SX4 Zxi BSIII	2.85	2007.0	Individual	110000	Petrol	Manual	15.00	15
1	Hyundai i20 Sportz 1.4 CRDi	4.70	2012.0	Dealer	70000	Diesel	Manual	21.90	13
2	Maruti Swift VDI BSIV	5.25	2015.0	Individual	70000	Diesel	Manual	25.20	12

Notice that Variable `full_name` contains both `Maker` and `Model` information - need to extract make and model from `full_name`.

```
df.shape
```

```
(13986, 11)
```

```
df.info();
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 13986 entries, 0 to 13985
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   full_name       13986 non-null  object
1   selling_price   13986 non-null  float64
2   year            13986 non-null  float64
3   seller_type     13986 non-null  object
4   km_driven       13986 non-null  int64
5   fuel_type       13986 non-null  object
6   transmission_type 13986 non-null  object
7   mileage         13986 non-null  float64
```

```

8   engine          13986 non-null float64
9   max_power       13986 non-null float64
10  seats           13986 non-null float64
dtypes: float64(6), int64(1), object(4)
memory usage: 1.2+ MB

```

```
display(df.describe())
```

	selling_price	year	km_driven	mileage	engine	max_power	seats
<b>count</b>	13986.000000	13986.000000	1.398600e+04	13986.000000	13986.000000	13986.000000	13986.000000
<b>mean</b>	7.377749	2014.516088	5.797629e+04	19.356221	1473.574905	97.685631	5.313242
<b>std</b>	9.368596	3.256795	5.378180e+04	4.638576	518.289204	45.067944	0.831653
<b>min</b>	0.325000	1991.000000	1.000000e+02	0.000000	0.000000	5.000000	2.000000
<b>25%</b>	3.400000	2013.000000	3.100000e+04	16.840000	1197.000000	73.900000	5.000000
<b>50%</b>	5.100000	2015.000000	5.120000e+04	19.160000	1248.000000	86.700000	5.000000
<b>75%</b>	7.800000	2017.000000	7.397800e+04	22.320000	1582.000000	112.000000	5.000000
<b>max</b>	395.000000	2021.000000	3.800000e+06	120.000000	6752.000000	626.000000	14.000000


- selling\_price ranges from [0.33, 395] - the values seem to be in lakhs.
- mileage min is 0, max is 120 - need to check for outliers.
- engine min is 0, max is 6752 - values appear to be in "cc" - need to check for outliers.

## ▼ Distributions

Lets check the distributions of these features one by one

```
df['selling_price'].hist(bins=50)
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f3a91998350>



Lets check the percentage of outliers. We are manually taking a call that price over a 100 (K) is an outlier by looking at the above plot

```
(df.loc[df.selling_price > 100].shape[0] / df.shape[0]) * 100
```

```
0.06435006435006435
```

Lets saturate selling price at **100**.

```
df.loc[df.selling_price > 100, 'selling_price'] = 100
```


- Can you think of an automated way of deciding this threshold 100 ? Seems like we made this decision arbitrarily
- **Hint:** use quantiles

## ▼ Outliers and Missing Values

Q: Which plot(s) is/are suitable for outlier visualisation ?

Similarly lets plot a histogram for **mileage**

```
df["mileage"].hist(bins=50);
```

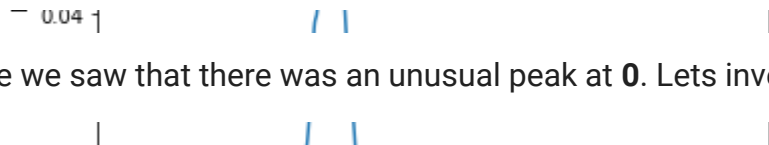


Another useful plot is called the 'density' plot. Which is also similar to histogram, but with a more *approximate* shape. It tries to fit a suitable distribution over the histogram and hence includes points that might not even be present in data at hand.

```
df["mileage"].plot.density();
```



We observe a skewed plot with right tail to long. An interesting point here would be mileage > 40, because after that it looks like the values might just be outliers.



Above we saw that there was an unusual peak at **0**. Lets investigate further. Looking at points where mileage is less than 5

```
df[df["mileage"]<=5].mileage.describe()
```

```
count    119.000000
mean      0.033613
std       0.366679
min       0.000000
25%       0.000000
50%       0.000000
75%       0.000000
max       4.000000
Name: mileage, dtype: float64
```

Looks like we just have a number of points with mileage 0. That does not make much sense as it cannot be a real value for car mileage. Investigating mileage=0 cases

```
df.loc[df.mileage == 0].shape[0] / df.shape[0] # 0.8 %
```

```
0.008437008437008437
```

```
df[df["mileage"] == 0.0].shape
```

```
(118, 11)
```

That's a very small number of points. We can safely impute, or drop these points later during pre-processing

### ▼ Are the top values erroneous?

Above we had also noticed that there was a significant point at mileage=40. After this there were only very few points and looked like outliers. Let's check them

```
df[df["mileage"] > 40].mileage.describe()
```

```
count      3.000000
mean      116.666667
std         5.773503
min       110.000000
25%       115.000000
50%       120.000000
75%       120.000000
max       120.000000
Name: mileage, dtype: float64
```

```
df[df["mileage"] > 40]
```

	full_name	selling_price	year	seller_type	km_driven	fuel_type	transmission_type	mileage
<b>4291</b>	Mahindra e2o Premium	3.25	2013.0	Individual	50000	Electric	Automatic	120.0
<b>4751</b>	Mahindra e2o T2	3.60	2015.0	Dealer	42862	Electric	Automatic	120.0

```
df[df["mileage"] > 40].fuel_type.value_counts()
```

```
Electric      3
Name: fuel_type, dtype: int64
```

- So all vehicles with mileage > 40 are electric vehicles which makes sense.


- So they are not incorrect (error) values.
- Let us not remove these points for now

## ▼ Correlations

Now lets look at the correlations of numerical features with each other

```
plt.figure(figsize=(10,8))  
ax = sns.heatmap(df.corr(), cmap="YlGnBu", annot=True)
```






We see the following correlations:

- Engine and max-power
- Max power and selling price Now based on our **domain knowledge** (about cars) we can digest these correlations. In this case, these correlations seem to make a lot of sense

Another visualisation we can check, is the scatter plot of each feature with the target. This can be easily done using pairplot from seaborn library



```
sns.pairplot(df, y_vars=["selling_price"]);
```



```
df['seats'].value_counts(normalize=True)
```

```
5.0    0.837123
7.0    0.115902
8.0    0.022237
4.0    0.011440
6.0    0.007293
9.0    0.003432
10.0   0.001502
2.0    0.001001
```

```
14.0      0.000072
Name: seats, dtype: float64
```

## ▼ Categorical Variables

### ▼ How can we extract Make and Model from "full\_name"

We have noticed that the name contains the information of the brand and the model. Let's extract those and create new features from them. After that we can drop the original column

```
df["make"] = df.full_name.apply(lambda x : x.split()[0])
df["model"] = df.full_name.apply(lambda x : " ".join(x.split()[1:]))
df = df.drop("full_name", axis=1)
df.head(2)
```

	selling_price	year	seller_type	km_driven	fuel_type	transmission_type	mileage	engine	max_po
0	2.85	2007.0	Individual	110000	Petrol	Manual	15.0	1586.0	104

Now let us look at some categorical variables one by one

```
display(df.describe(include="object"))
```

**seller\_type fuel\_type transmission\_type make model**

```
df["fuel_type"].value_counts()
```

```
Petrol      6870
Diesel      6823
CNG         233
LPG         49
Electric    11
Name: fuel_type, dtype: int64
```

```
df["transmission_type"].value_counts()
```


```
Manual      11251
Automatic   2735
Name: transmission_type, dtype: int64
```

```
df["seller_type"].value_counts()
```

```
Dealer      8393
Individual  5450
Trustmark Dealer  143
Name: seller_type, dtype: int64
```

### ▼ Does selling price vary in some of these groups

```
df.loc[df.selling_price<=df.selling_price.quantile(0.9)].groupby('seller_type')['selling_price'].plot.density();
plt.legend();
```



Looks like for different seller type, the selling price distribution is not very different. Hence one intuition that can be developed from this plot is that the **seller\_type** variable will not turn out to be very important for the ML model. (We will check this later)



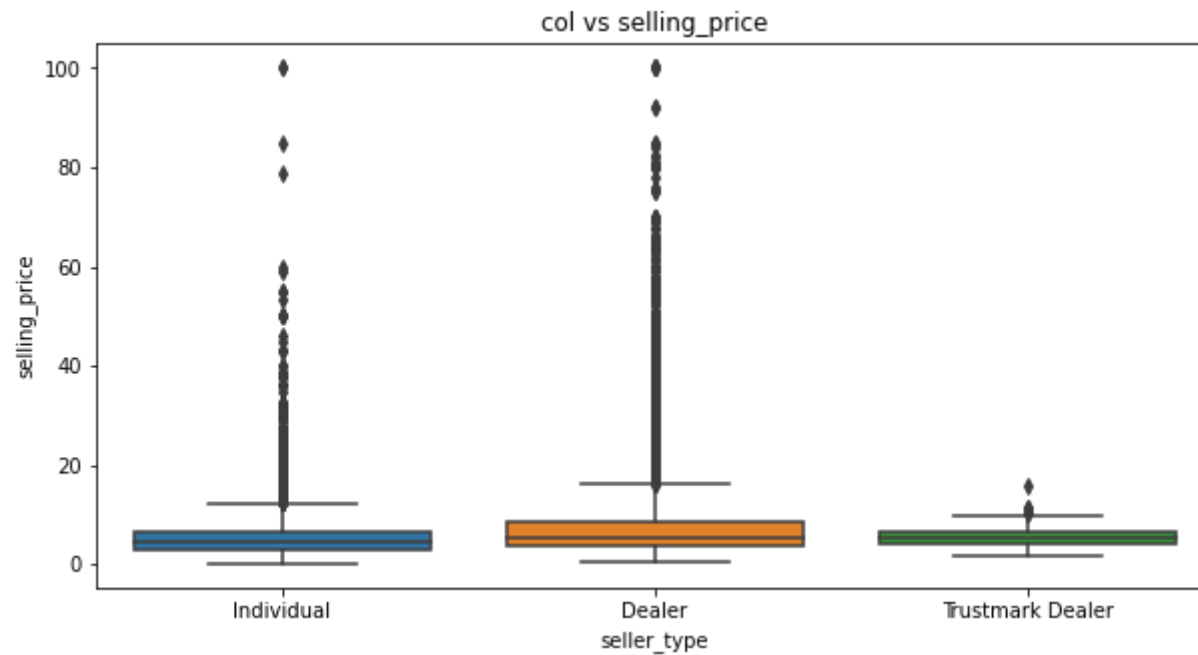
0.000 | ————— |

### ▼ Lets check for all groups?

Another way to visualise the same is to use a plot known as **box-plot**

Note: Although, `seats` is numeric, but we will consider it as an ordinal variable and perform similar analysis as we do for other categorical variables.

```
for col in ['seller_type', 'fuel_type', 'seats']:
    plt.figure(figsize=(10,5))
    sns.boxplot(y='selling_price',x=col, data=df)
    plt.title('col vs selling_price')
    plt.show()
```




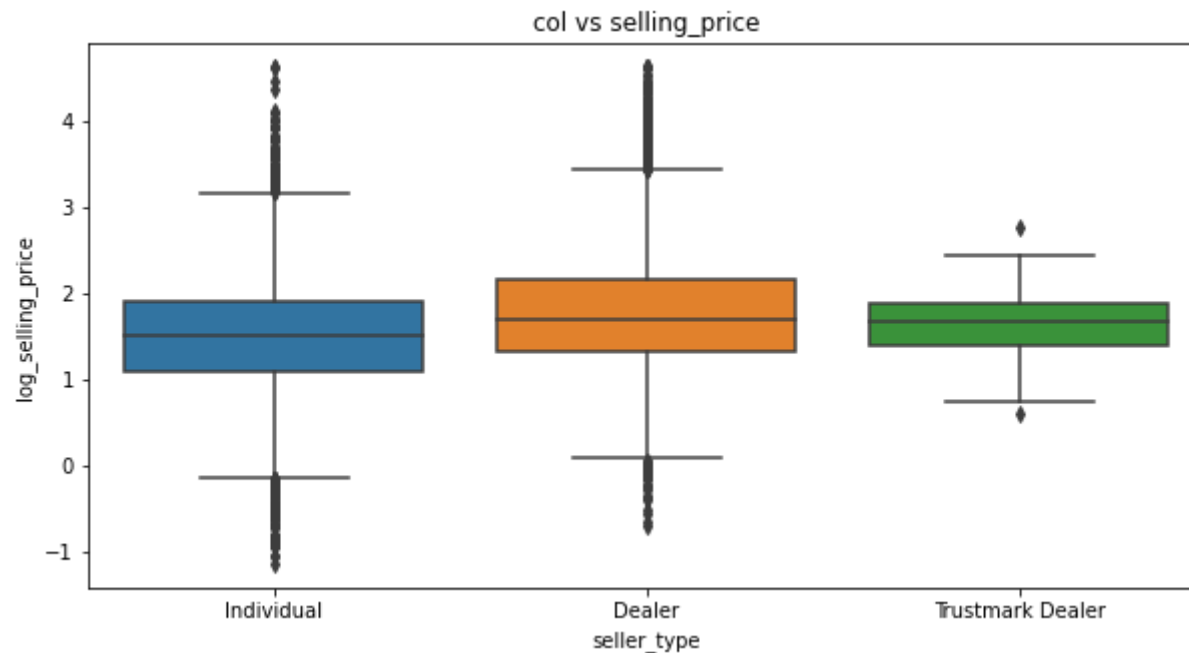


▼ How can we transform the scale of selling-price for better visualisation?



```
df["log_selling_price"] = np.log(df["selling_price"].values)
```

```
for col in ['seller_type', 'fuel_type', 'seats']:
    plt.figure(figsize=(10,5))
    sns.boxplot(y='log_selling_price',x=col, data=df)
    plt.title('col vs selling_price')
    plt.show()
```



- selling\_price of "coupe" is quite variational - important to be kept for predictive modelling

### ▼ How many categories of "make" in the data?

```
len(df["make"].unique())
```

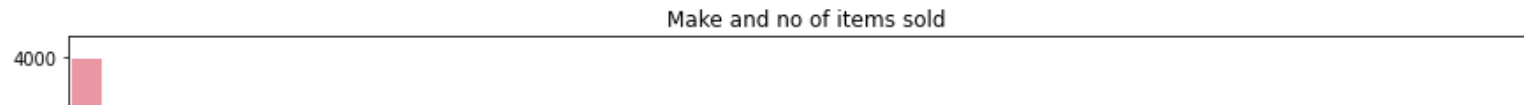
41

There are 41 car brands in the data - too many levels, think of possible ways to encode them as numbers

### ▼ What are item counts for different "make"?

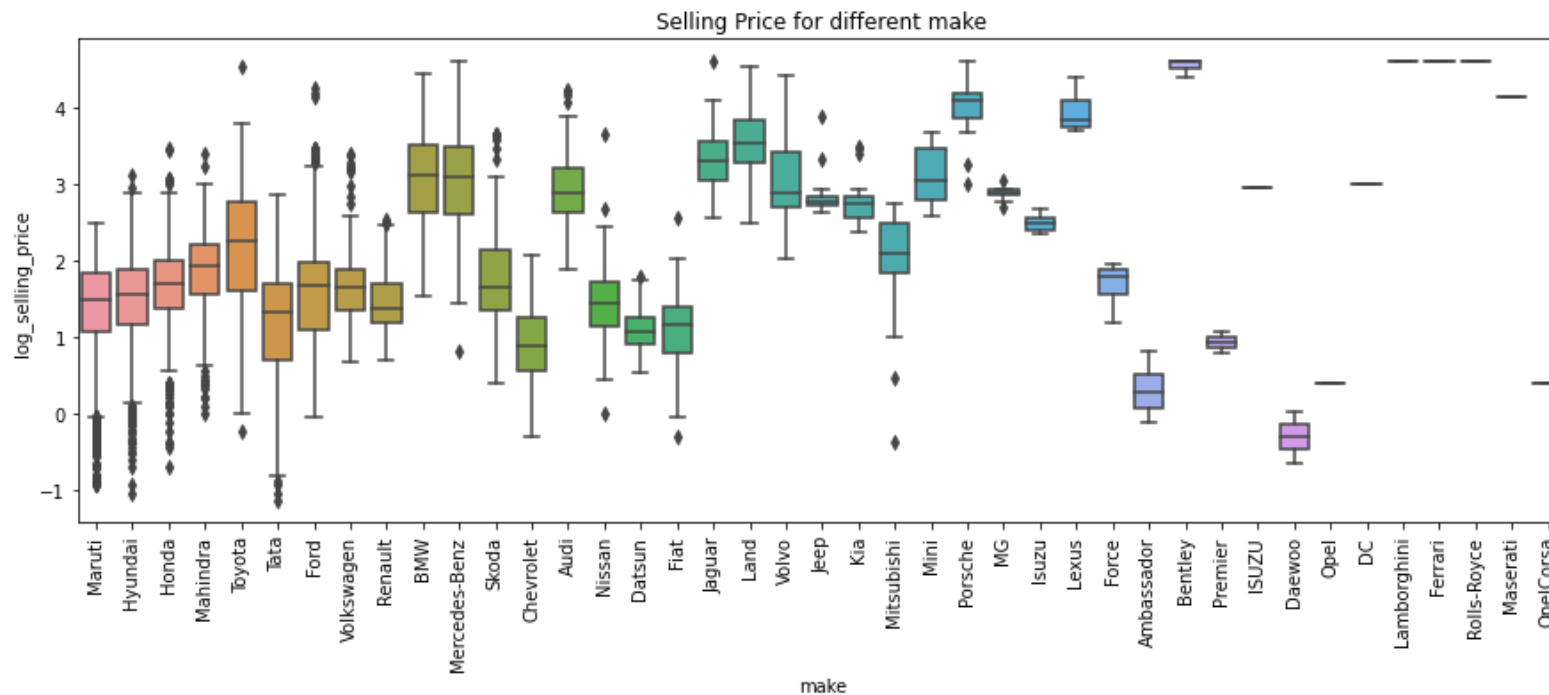
```
plt.figure(figsize=(15,5))
sns.countplot(x='make', data=df, order = df['make'].value_counts().index)
plt.xticks(rotation=90);
plt.title("Make and no of items sold")
plt.show()
```





▼ How is "make" related to the selling price?

```
plt.figure(figsize=(15,5))
sns.boxplot(y='log_selling_price',x='make', data=df, order = df['make'].value_counts().index)
plt.xticks(rotation=90);
plt.title("Selling Price for different make")
plt.show()
```



Lets drop this column for now

```
df.drop("log_selling_price", axis=1, inplace=True)
```

- How should we treat number of seats? Continuous OR Categorical ?
- **Answer:**
  - Number of seats is a variable with a finite discrete set of values. However it is ordered. If we make it categorical, it would be fed into the model as unordered.
  - But we need to think do we want that order? Example 2 seater cars are most likely more expensive than other cars. So in this sense  $2 > 5$ . We actually would benefit from discarding the order.
- Can you convert number of seats to a variable which has 9 unique values with 3 categories?
- **Hint:** Merge some elements. Eg all cars with  $> 5$  seats are on category.
- Can we do something similar with the "make variable"

## ▼ Feature Engineering

```
df.head()
```

	selling_price	year	seller_type	km_driven	fuel_type	transmission_type	mileage	engine	max_po
0	2.85	2007.0	Individual	110000	Petrol	Manual	15.00	1586.0	104
1	4.70	2012.0	Dealer	70000	Diesel	Manual	21.90	1396.0	86
2	5.25	2015.0	Individual	70000	Diesel	Manual	25.20	1248.0	74
...	...	...	...	...	...	...	...	...	...


## ▼ 1. Year

Sometimes it is good practice to convert the features into the way we actually understand. This will not necessarily help the model, but helps in us in understanding / interpreting predictions little better

- So lets convert year to more readable **age**

```
current_year = 2020  
age = current_year - df['year']  
plt.scatter(age, df.selling_price)
```

<matplotlib.collections.PathCollection at 0x7f3a8eafdb10>



**Higher selling price for newer cars !!**

## ▼ 2. Categorical Encoding

Q: How do we encode categorical variables for modelling?

2.1 transmission\_type (2 categories)

```
# explains how two category variable can be encoded as a binary variable
transmission_dummy = pd.get_dummies(df["transmission_type"])
df = pd.concat([df, transmission_dummy], axis=1)
df = df.drop(["transmission_type", "Automatic"], axis=1)
df = df.rename(columns={"Manual": "transmission_type"})
df.sample(frac=1).head(2)
```

	selling_price	year	seller_type	km_driven	fuel_type	mileage	engine	max_power	seats	mak
<b>4755</b>	7.75	2017.0	Dealer	114000	Diesel	24.3	1248.0	88.5	5.0	Maru

## ▼ 2.2 fuel\_type (5 categories)

```
fuel_type_dummy = pd.get_dummies(df["fuel_type"], prefix="seats")
fuel_type_dummy.sample(frac=1).head(2)
```

	seats_CNG	seats_Diesel	seats_Electric	seats_LPG	seats_Petrol
<b>2013</b>	0	0	0	0	1
<b>9172</b>	0	0	0	0	1

**NOTE:** We remove 1 of the dummy variables to avoid the dummy variable trap

```
fuel_type_dummy = fuel_type_dummy[fuel_type_dummy.columns[1:]]
fuel_type_dummy.head(2)
```

**seats\_Diesel seats\_Electric seats\_LPG seats\_Petrol**

```
pd.concat([df, fuel_type_dummy], axis=1).sample(frac=1).head(2)
```

	selling_price	year	seller_type	km_driven	fuel_type	mileage	engine	max_power	seats	make
<b>3564</b>	18.99	2018.0	Dealer	20306	Diesel	15.10	2179.0	152.87	7.0	Mahindra
<b>989</b>	7.50	2016.0	Dealer	62563	Diesel	28.09	1248.0	88.50	5.0	Maruti

## ▼ 2.3 make (too many categories)

```
df.make.nunique()
```

```
41
```

- How to handle so many categories?
- **Hint:** Here we have too many categories. We can merge categories as in 'seats' variable or we can try target encoding. Dummy encoding would make too many features with highly diluted information.

Or we can try another technique known as **Target Encoding**

```
encoded_make = df.groupby('make')['selling_price'].transform('mean')
```

```
show_df = df[['make']].copy()
```

```
show_df['encoded_make'] = encoded_make  
show_df = show_df.drop_duplicates()  
print(len(show_df))  
show_df.head()
```

41

	make	encoded_make
0	Maruti	4.648484
1	Hyundai	5.344764
3	Honda	5.911203
4	Volkswagen	5.632862
5	Toyota	11.179478

### ▼ 3. Scaling

Q: Now, All these variables have different ranges. Does that work?

```
df.describe()
```

	selling_price	year	km_driven	mileage	engine	max_power	seats
<b>count</b>	13986.000000	13986.000000	1.398600e+04	13986.000000	13986.000000	13986.000000	13986.000000
<b>mean</b>	7.327126	2014.516088	5.797629e+04	19.356221	1473.574905	97.685631	5.313242

Lets scale this data. We will use MinMaxScaler here (in practice you can choose your scaling techniques depending on the problem, or just try all and see what works best)

```
from sklearn.preprocessing import MinMaxScaler, StandardScaler
```

```
scaler = MinMaxScaler()
scaler.fit(df[['selling_price', 'km_driven', 'mileage']])
scaled_values = scaler.transform(df[['selling_price', 'km_driven', 'mileage']]) # returns numpy.ndarray not df.
scaled_df = pd.DataFrame(scaled_values, columns=['selling_price', 'km_driven', 'mileage'])
scaled_df.head()
```

	selling_price	km_driven	mileage
<b>0</b>	0.025332	0.028922	0.125000
<b>1</b>	0.043893	0.018395	0.182500
<b>2</b>	0.049411	0.018395	0.210000
<b>3</b>	0.009280	0.023659	0.108333
<b>4</b>	0.043391	0.010763	0.137250

```
scaled_df.describe()
```

	selling_price	km_driven	mileage
<b>count</b>	13986.000000	13986.000000	13986.000000
<b>mean</b>	0.070250	0.015231	0.161302
<b>std</b>	0.083378	0.014153	0.038655
<b>min</b>	0.000000	0.000000	0.000000
<b>25%</b>	0.020850	0.008133	0.110333
<b>50%</b>	0.030850	0.008133	0.110333
<b>75%</b>	0.074994	0.019442	0.186000

- Try standard scaler and print "df.describe", and make density plots for both std scaler and minmax scaler

## ▼ Sticking it all together

```
def merge_seats(x):
    if 2 <= x <= 4:
        return '2-4'
    elif x > 5:
        return '>5'
    else:
        return '5'

def preprocess(df):
    df = df.loc[df.mileage != 0].copy()
    outlier_theshold = df.selling_price.quantile(0.95)
    df.loc[df.selling_price > outlier_theshold, 'selling_price'] = outlier_theshold
    df['age'] = pd.to_datetime('now', utc=True).date().year - df['year']
    df['full_name'] = df['full_name'].str.upper()
    df['seats'] = df['seats'].apply(merge_seats)
    df.drop(columns=['year'])
    return df

def feature_engineering(df):
```



```

df['make'] = df.full_name.apply(lambda x : x.split()[0])
df['model'] = df.full_name.apply(lambda x : " ".join(x.split()[1:]))
one_hot_encode_cols = ['seller_type', 'fuel_type', 'transmission_type', 'seats']
for column in one_hot_encode_cols:
    df = pd.concat([df, pd.get_dummies(df[column]).iloc[:, 1:]], axis=1)
df.drop(columns=one_hot_encode_cols + ['full_name'], inplace=True)
df['make'] = df.groupby('make')['selling_price'].transform('median')
df['model'] = df.groupby('model')['selling_price'].transform('median')
scaler = MinMaxScaler()
df = pd.DataFrame(scaler.fit_transform(df), columns=df.columns)
return df

```

```

df = pd.read_csv("train-cars24-car-price.csv")
df = preprocess(df)
df = feature_engineering(df)
df.head()

```

	selling_price	year	km_driven	mileage	engine	max_power	age	make	model	Indiv
0	0.122128	0.517241	0.028922	0.094828	0.234893	0.160515	0.482759	0.179031	0.093108	
1	0.211608	0.689655	0.018395	0.154310	0.206754	0.134879	0.310345	0.203759	0.200726	
2	0.238210	0.793103	0.018395	0.182759	0.184834	0.111111	0.206897	0.179031	0.224184	
3	0.044740	0.448276	0.023659	0.077586	0.198904	0.136876	0.551724	0.233432	0.087062	
4	0.209190	0.793103	0.010763	0.107500	0.177429	0.111111	0.206897	0.218595	0.229504	

```
df.shape
```


```
(13868, 18)
```

## ▼ Correlations with all numeric features

And before we wrap up, lets just see correlations of all variables since now we have all usable information in numerical form

```
plt.figure(figsize=(15, 15))  
sns.heatmap(df.corr(), cmap="YlGnBu", annot=True)
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f3a8eae9690>



- max\_power, engine and transmission\_type are highest correlated with selling\_price in same order.
- New encoded features are weak predictor, but they together would might make strong predictors.



